

```
In [53]: import pandas as pd
```

```
In [54]: data=pd.read_csv("/home/placement/Downloads/fiat500.csv")
```

```
In [55]: data.describe()
```

```
Out[55]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>count</b>	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
<b>mean</b>	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
<b>std</b>	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
<b>min</b>	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
<b>25%</b>	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
<b>50%</b>	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
<b>75%</b>	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
<b>max</b>	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

```
In [56]: data.head()
```

```
Out[56]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>0</b>	1	lounge	51	882	25000	1	44.907242	8.611560	8900
<b>1</b>	2	pop	51	1186	32500	1	45.666359	12.241890	8800
<b>2</b>	3	sport	74	4658	142228	1	45.503300	11.417840	4200
<b>3</b>	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
<b>4</b>	5	pop	73	3074	106880	1	41.903221	12.495650	5700

```
In [57]: data1=data.drop(['lat','lon','ID'],axis=1)#unwanted columns removed  
data1
```

Out[57]:

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [58]: data=pd.get_dummies(data)
data
```

Out[58]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price	model_lounge	model_pop	model_sport
0	1	51	882	25000	1	44.907242	8.611560	8900	1	0	0
1	2	51	1186	32500	1	45.666359	12.241890	8800	0	1	0
2	3	74	4658	142228	1	45.503300	11.417840	4200	0	0	1
3	4	51	2739	160000	1	40.633171	17.634609	6000	1	0	0
4	5	73	3074	106880	1	41.903221	12.495650	5700	0	1	0
...	...	...	...	...	...	...	...	...	...	...	...
1533	1534	51	3712	115280	1	45.069679	7.704920	5200	0	0	1
1534	1535	74	3835	112000	1	45.845692	8.666870	4600	1	0	0
1535	1536	51	2223	60457	1	45.481541	9.413480	7500	0	1	0
1536	1537	51	2557	80750	1	45.000702	7.682270	5990	1	0	0
1537	1538	51	1766	54276	1	40.323410	17.568270	7900	0	1	0

1538 rows × 11 columns

```
In [59]: data.shape
```

Out[59]: (1538, 11)

```
In [60]: data2=pd.get_dummies(data1)
```

```
In [61]: data2
```

```
Out[61]:
```

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...	...	...	...	...	...	...	...	...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

```
In [62]: data2.shape
```

```
Out[62]: (1538, 8)
```

```
In [63]: y=data2['price']  
x=data2.drop('price',axis=1)  
#predicted value we removed from data frame
```

In [64]:

y

Out[64]:

```
0      8900
1      8800
2      4200
3      6000
4      5700
```

```
...
1533   5200
1534   4600
1535   7500
1536   5990
1537   7900
```

Name: price, Length: 1538, dtype: int64

In [65]:

```
#!pip install scikit-learn
```

In [66]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
#divide the columns into testing and training
```

In [67]:

```
x_test.head(5)
```

Out[67]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0

```
In [68]: y_test.head(5)
```

```
Out[68]: 481      7900
         76      7900
         1502     9400
         669     8500
         1409     9700
         Name: price, dtype: int64
```

```
In [69]: x_train.head(5)
```

```
Out[69]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
527	51	425	13111	1	1	0	0
129	51	1127	21400	1	1	0	0
602	51	2039	57039	1	0	1	0
331	51	1155	40700	1	1	0	0
323	51	425	16783	1	1	0	0

```
In [70]: y_train.head(5)
```

```
Out[70]: 527      9990
         129      9500
         602      7590
         331      8750
         323      9100
         Name: price, dtype: int64
```

```
In [71]: from sklearn.linear_model import LinearRegression
         reg=LinearRegression() #creating object of LinearRegression
         reg.fit(x_train,y_train) #training and fitting LR object using training data
```

```
Out[71]: LinearRegression()
```

```
In [72]: y_pred=reg.predict(x_test)
         #prediction price
```

```
In [73]: y_pred
8840.08397206, 9916.27565791, 10287.45603992, 9964.3213269 ,
8403.51255128, 9345.81907605, 8521.46225147, 9743.68712672,
9791.34520178, 9779.16293972, 6753.27416058, 7354.16762745,
8760.24542762, 9923.66596418, 9812.92276721, 10466.90125415,
8163.46726237, 6659.46839415, 9987.65677522, 8866.7826029 ,
9952.37340054, 10187.72427693, 10231.39378767, 10091.11325493,
9365.98570732, 10009.10088406, 9141.00566394, 10099.11667176,
7803.77049829, 6009.84398185, 8800.33824151, 10237.60733785,
5609.98366311, 10097.61555355, 9684.99946572, 7644.67379732,
9276.37891542, 7371.5492091 , 10287.98873148, 10067.26428381,
10552.64805598, 9966.72383894, 10068.46126756, 6232.53552963,
10584.55044373, 9965.98687522, 10529.44404458, 9602.67646085,
9665.77720284, 6186.06948587, 8073.87436253, 10345.58323918,
6344.74803956, 7361.62678204, 10058.57116223, 6792.219309 ,
7897.72464823, 5261.45936067, 4540.24137423, 8709.36468047,
6882.0117409 , 7406.73353952, 6795.61189392, 7047.27998963,
9945.33400083, 8856.93910595, 9378.02074127, 10389.561154 ,
10092.46332921, 10381.52000388, 9723.92466625, 5996.3331428 ,
9786.14866981, 7708.49649098, 5583.48163469, 4932.92788329,
9856.66053994, 9236.22981005, 10092.64052142, 6256.43516278,
```

```
In [74]: from sklearn.metrics import r2_score
         r2_score(y_test,y_pred)
```

```
Out[74]: 0.8415526986865394
```

```
In [75]: from sklearn.metrics import mean_squared_error
         mean_squared_error(y_pred,y_test)
```

```
Out[75]: 581887.727391353
```

```
In [76]: import math  
print(math.sqrt(581887.727391353))
```

762.8156575420782

```
In [77]: y_pred
```

```
8456.30006203, 6499.76668237, 7768.57829985, 6832.86406122,  
8347.96113362, 10439.02404036, 7356.43463051, 8562.56562053,  
9820.78555199, 10035.83571539, 7370.77198022, 9411.45894006,  
10352.85155564, 8045.21588007, 10446.80664758, 3736.20118868,  
10348.63930496, 10435.96627494, 6167.80169017, 10390.11317804,  
6527.69471073, 9116.4755691, 10484.52829, 9335.69889855,  
6709.57413543, 3390.72353093, 10106.33753331, 9792.46732008,  
6239.49568346, 4996.26346266, 9044.38667681, 9868.09959448,  
5484.13199252, 5698.5954821, 10086.86206874, 8115.81693479,  
10392.37800936, 6835.6573351, 6657.61744836, 5738.50576764,  
8896.80120764, 9952.37340054, 10390.28377419, 9419.10788866,  
9082.56591129, 10122.82465116, 10410.00504522, 10151.77663915,  
9714.85367238, 9291.92963633, 10346.99073888, 5384.22311343,  
9772.85146492, 6069.77107828, 9023.26394782, 10220.56195956,  
  
9238.89392583, 9931.47195375, 8321.42715662, 8377.80491069,  
7528.53327408, 10552.64805598, 10465.02437243, 10110.68940664,  
10238.17869436, 6841.77264488, 9625.64505547, 10412.59988875,  
9653.06224923, 7948.63618724, 9704.82523573, 7971.05970955,  
10399.51752022, 9176.43567301, 5803.03205787, 6698.19524313,
```



```
In [78]: Results=pd.DataFrame(columns=['price','predicted'])
Results['price']=y_test
Results['predicted']=y_pred
Results=Results.reset_index()
Results['Id']=Results.index
Results.head(15)
```

Out[78]:

	index	price	predicted	Id
0	481	7900	5867.650338	0
1	76	7900	7133.701423	1
2	1502	9400	9866.357762	2
3	669	8500	9723.288745	3
4	1409	9700	10039.591012	4
5	1414	9900	9654.075826	5
6	1089	9900	9673.145630	6
7	1507	9950	10118.707281	7
8	970	10700	9903.859527	8
9	1198	8999	9351.558284	9
10	1088	9890	10434.349636	10
11	576	7990	7732.262557	11
12	965	7380	7698.672401	12
13	1488	6800	6565.952404	13
14	1432	8900	9662.901035	14

```
In [79]: Results['diff']=Results.apply(lambda row:row.price-row.predicted,axis=1)
```

In [80]: Results

Out[80]:

	index	price	predicted	ld	diff
<b>0</b>	481	7900	5867.650338	0	2032.349662
<b>1</b>	76	7900	7133.701423	1	766.298577
<b>2</b>	1502	9400	9866.357762	2	-466.357762
<b>3</b>	669	8500	9723.288745	3	-1223.288745
<b>4</b>	1409	9700	10039.591012	4	-339.591012
...	...	...	...	...	...
<b>503</b>	291	10900	10032.665135	503	867.334865
<b>504</b>	596	5699	6281.536277	504	-582.536277
<b>505</b>	1489	9500	9986.327508	505	-486.327508
<b>506</b>	1436	6990	8381.517020	506	-1391.517020
<b>507</b>	575	10900	10371.142553	507	528.857447

508 rows × 5 columns

```
In [81]: Results=pd.DataFrame(columns=['price', 'predicted'])  
Results['price']=y_test  
Results['predicted']=y_pred  
Results.head(15)
```

Out[81]:

	price	predicted
481	7900	5867.650338
76	7900	7133.701423
1502	9400	9866.357762
669	8500	9723.288745
1409	9700	10039.591012
1414	9900	9654.075826
1089	9900	9673.145630
1507	9950	10118.707281
970	10700	9903.859527
1198	8999	9351.558284
1088	9890	10434.349636
576	7990	7732.262557
965	7380	7698.672401
1488	6800	6565.952404
1432	8900	9662.901035

```
In [82]: Results['diff']=Results.apply(lambda row:row.price-row.predicted,axis=1)
```

In [83]: Results

Out[83]:

	price	predicted	diff
<b>481</b>	7900	5867.650338	2032.349662
<b>76</b>	7900	7133.701423	766.298577
<b>1502</b>	9400	9866.357762	-466.357762
<b>669</b>	8500	9723.288745	-1223.288745
<b>1409</b>	9700	10039.591012	-339.591012
...	...	...	...
<b>291</b>	10900	10032.665135	867.334865
<b>596</b>	5699	6281.536277	-582.536277
<b>1489</b>	9500	9986.327508	-486.327508
<b>1436</b>	6990	8381.517020	-1391.517020
<b>575</b>	10900	10371.142553	528.857447

508 rows × 3 columns

In [ ]: